



GDG DevFest Bucharest 2018



Architecting for the Google Cloud Platform



Márton Kodok / [@martонkодок](https://twitter.com/martонkодок)
Google Developer Expert at REEA.net - Targu Mures

November 2018 - Bucharest, Romania

About me

- Geek. Hiker. Do-er.
- Among the Top3 romanians on Stackoverflow 127k reputation
- Google Developer Expert on Cloud technologies
- Crafting Web/Mobile backends at **REEA.net**
- BigQuery/Redis and database engine expert
- Active in mentoring and IT community

Slideshare: **martonkodok**
Twitter: **@martonkodok**
StackOverflow: **pentium10**
GitHub: **pentium10**

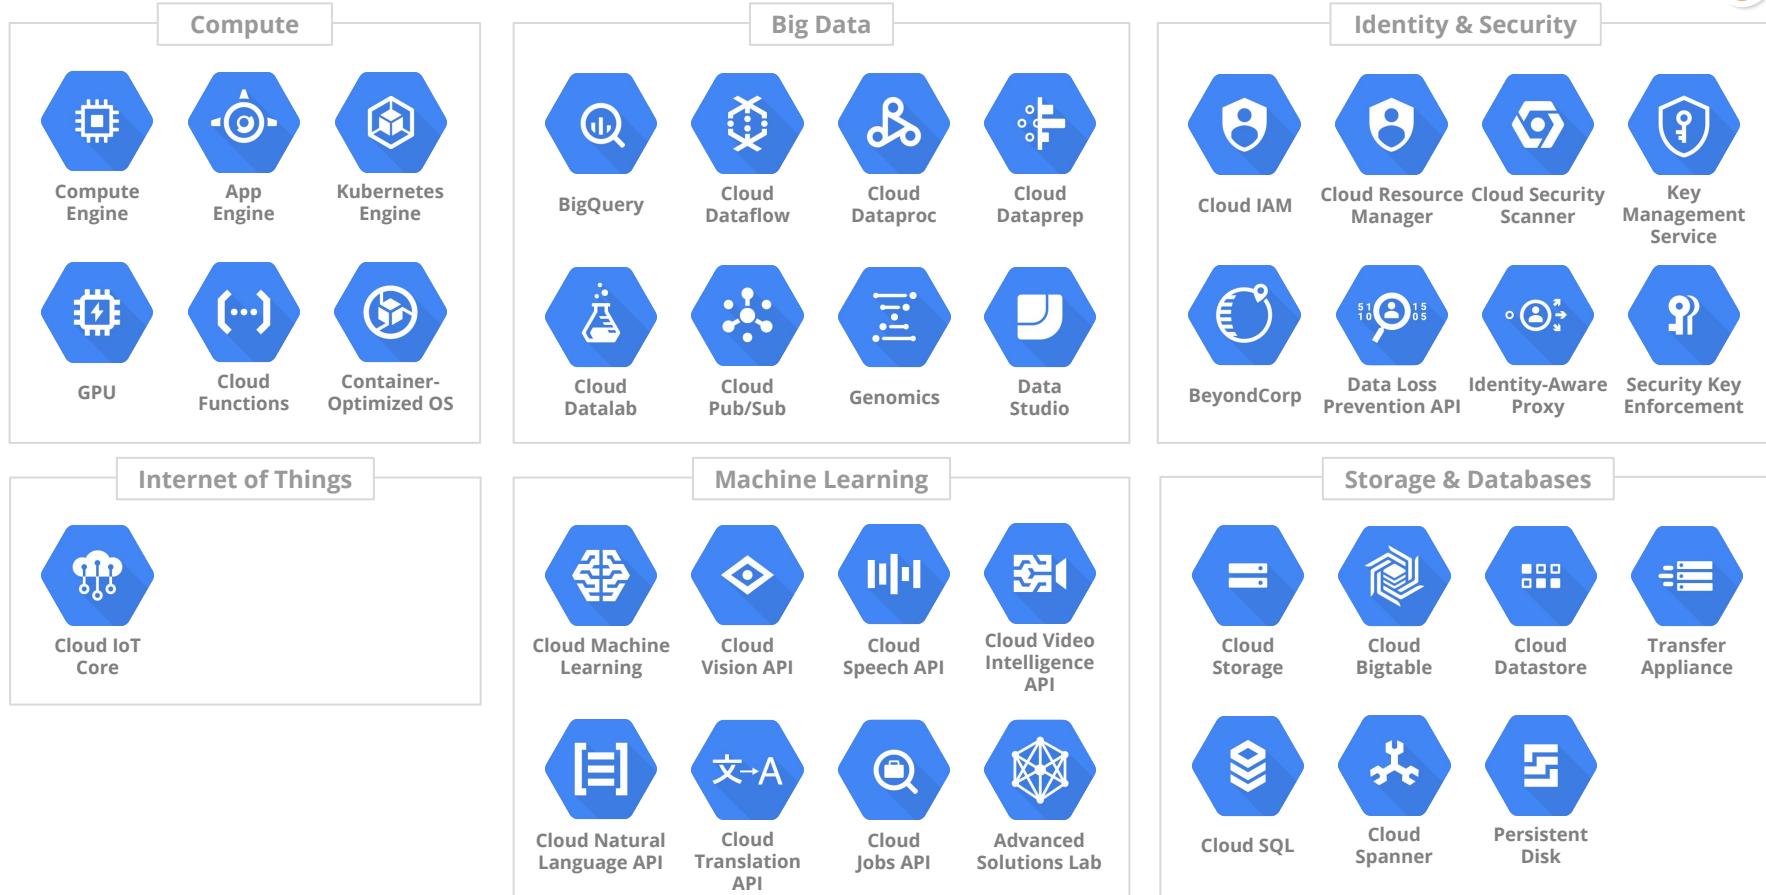


Build on the same infrastructure
that powers Google



Google Cloud Platform

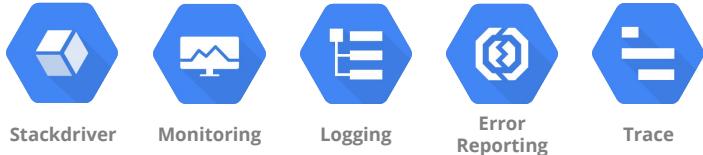
Google Cloud Platform (GCP)



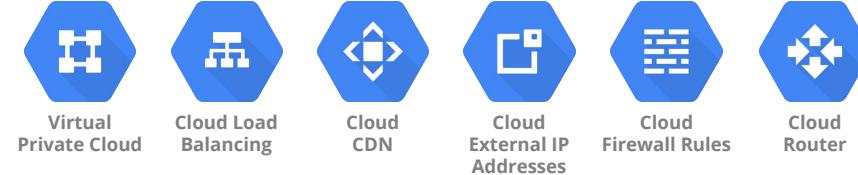
Google Cloud Platform (GCP)



Management Tools



Networking



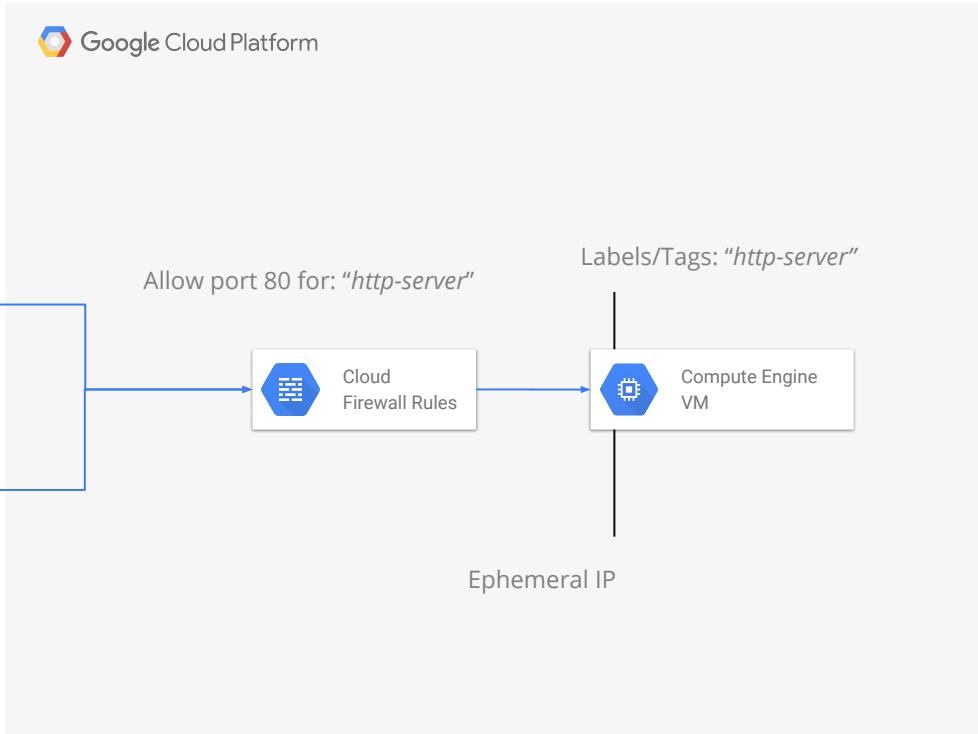
Developer Tools



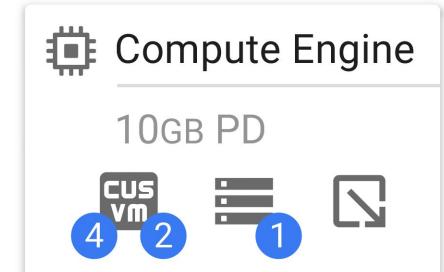
Agenda

1. Virtual Machines - Compute Engine
2. **Autoscaling** - Design for resiliency, scalability, and disaster recovery
3. Database layer - Identification of storage needs and mapping
4. **Serverless** - App Engine platform - Cloud Functions
5. Microservices architecture
6. **Task Queues** - Cloud Pub/Sub - DataFlow
7. Data warehouse in the Cloud
8. Practical use cases
9. Qwiklabs + Coursera

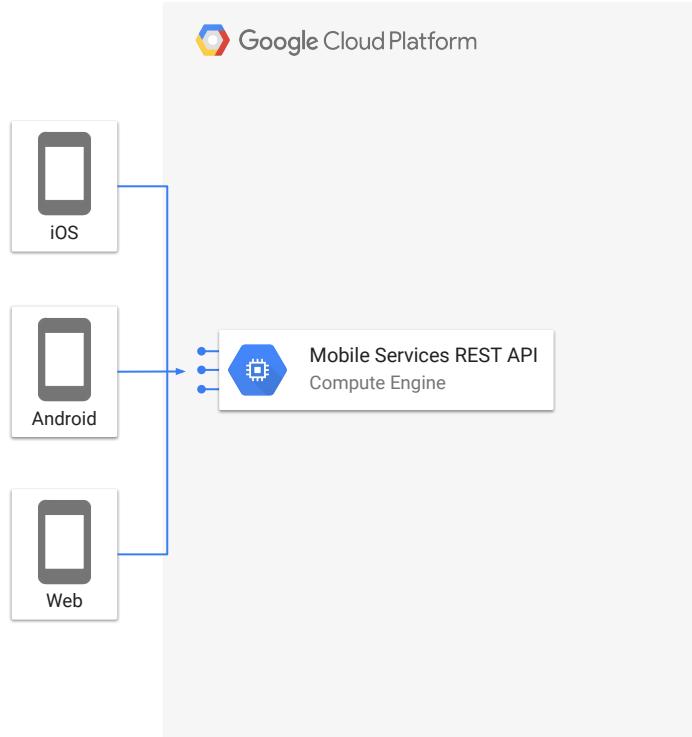
Compute Engine - Virtual Machine in the Cloud



VM details



Compute Engine - General Use



Flexibility for Every Workload

- High-Performance, Scalable VMs
- Custom Machine Types
- Persistent Disks or SSD (local or network based)
- Per-second billing
- Networking
- Firewall
- Load balancing
- Compliance & Security

Creating a VM instance steps



Creating a VM instance

Create an instance

Name [?](#)
frontend-webtier

Region [?](#) Zone [?](#)
us-east1 (South Carolina) us-east1-b

Machine type
Customize to select cores, memory and GPUs.
1 vCPU 3.75 GB memory Customize

Container [?](#)
 Deploy a container image to this VM instance. [Learn more](#)

Boot disk [?](#)
New 10 GB standard persistent disk
Image: Debian GNU/Linux 9 (stretch) Change

Identity and API access [?](#)
Service account [?](#) Compute Engine default service account

Access scopes [?](#)

\$24.67 per month estimated
Effective hourly rate \$0.034 (730 hours per month)

Item	Estimated costs
1 vCPU + 3.75 GB memory	\$34.67/month
10 GB standard persistent disk	\$0.40/month
Sustained use discount ?	- \$10.40/month
Total	\$24.67/month

[Compute engine pricing](#) ↗

▲ Less





Creating a VM instance - Choosing Region, Instance type

Create an instance

- asia-east1 (Taiwan)
- asia-northeast1 (Tokyo)
- asia-south1 (Mumbai)
- asia-southeast1 (Singapore)
- australia-southeast1 (Sydney)
- europe-north1 (Finland)
- europe-west1 (Belgium)
- europe-west2 (London)
- europe-west3 (Frankfurt)
- europe-west4 (Netherlands)
- northamerica-northeast1 (Montréal)
- southamerica-east1 (São Paulo)
- us-central1 (Iowa)
- us-east1 (South Carolina) Default
- us-east4 (Northern Virginia)
- us-west1 (Oregon)

Zone

GB memory

Identity and API access

Service account

Create an instance

4 vCPUs

- micro (1 shared vCPU)
0.6 GB memory, f1-micro
- small (1 shared vCPU)
1.7 GB memory, g1-small
- 1 vCPU
3.75 GB memory, n1-standard-1
- 2 vCPUs
7.5 GB memory, n1-standard-2
- 4 vCPUs
15 GB memory, n1-standard-4
- 8 vCPUs
30 GB memory, n1-standard-8
- 16 vCPUs
60 GB memory, n1-standard-16
- 32 vCPUs
120 GB memory, n1-standard-32
- 64 vCPUs
240 GB memory, n1-standard-64
- 96 vCPUs
360 GB memory, n1-standard-96

Firewall





Creating a VM instance - OS, boot disk, Network

Boot disk

Select an image or snapshot to create a boot disk; or attach an existing disk

[OS images](#) [Application images](#) [Custom images](#) [Snapshots](#) [Existing disks](#)

- Centos 7
x86_64 built on 20180611
- CoreOS alpha 1800.1.0
amd64-user published on 2018-06-11
- CoreOS beta 1772.4.0
amd64-user published on 2018-06-14
- CoreOS stable 1745.7.0
amd64-user published on 2018-06-14
- Ubuntu 14.04 LTS
amd64 trusty image built on 2018-06-14
- Ubuntu 16.04 LTS
amd64 xenial image built on 2018-06-12
- Ubuntu 17.10
amd64 artful image built on 2018-06-12
- Ubuntu 18.04 LTS
amd64 bionic image built on 2018-06-13
- Container-Optimized OS 68-10718.23.0 beta
Kernel: ChromiumOS-4.14.41 Kubernetes: 1.10.2 Docker: 17.03.2
- Container-Optimized OS 69-10783.0.0 dev
Kernel: ChromiumOS-4.14.48 Kubernetes: 1.10.4 Docker: 17.03.2

Can't find what you're looking for? Explore hundreds of VM solutions in [Cloud Launcher](#)

Boot disk type [?](#) Size (GB) [?](#)

SSD persistent disk

10

Create an instance

Identity and API access [?](#)

Service account [?](#)

Compute Engine default service account

Access scopes [?](#)

- Allow default access
- Allow full access to all Cloud APIs
- Set access for each API

Firewall [?](#)

Add tags and firewall rules to allow specific network traffic from the Internet

Allow HTTP traffic

Allow HTTPS traffic

Management

Disks

Networking

SSH Keys

Network tags [?](#) (Optional)

Network interfaces [?](#)

default default (10.142.0.0/20)





VM instances

Google Cloud Platform marton-data

VM instances

CREATE INSTANCE

VM instances

Filter VM instances

Name	Zone	Recommendation	Internal IP	External IP	Connect
instance-1	us-west1-c		10.138.0.2 (nic0)	None	SSH

Compute Engine

VM instances

Instance groups

Instance templates

Disks

Snapshots

Images





Virtual Private Network - Firewall rules

VPC network	Firewall rules						
	Ingress		Egress				
	Name	Targets	Source filters	Protocols / ports	Action	Priority	Network ^
VPC networks	<input type="checkbox"/> apache	Apply to all	IP ranges: 0.0.0.0/0	tcp:80,443	Allow	1000	default
External IP addresses	<input type="checkbox"/> beanstalkd	Apply to all	IP ranges: 173.193.140.142	tcp:11300	Allow	1000	default
Firewall rules	<input type="checkbox"/> cerebro-for-elasticsearch	Apply to all	IP ranges: 172.243.230.13, 1 more ▾	tcp:9000,5601	Allow	1000	default
Routes	<input type="checkbox"/> .elasticsearch	Apply to all	IP ranges: 5.15.22.26	tcp:9200	Allow	1000	default
VPC network peering	<input type="checkbox"/> .elasticsearch-cluster-2-node-0-firewall	elasticsearch-cluster-2-node-0	IP ranges: 0.0.0.0/0	tcp:22	Allow	1000	default
Shared VPC							





Cost savings! “Preemptible” - spot instances

- Much lower price than normal instances
- For fault-tolerant applications
- Can withstand possible instance losses
- Runs for maximum 24 hours
- Might not be always available
- Instance templates

Good for: Nodes (eg: Elasticsearch)
Batch jobs
Worker async tasks

Pricing: ranges 3-4 times lower price





Autoscaling with Instance Groups





Creating an Instance Group

VM instances

Instance groups

Instance templates

Disks

Snapshots

Images

TPUs

Committed use discounts

Metadata

Cloud Launcher

<

Location
Multi-zone groups span multiple zones which assures higher availability
[Learn more](#)

Single-zone
 Multi-zone

Region [?](#) us-central1 (Iowa)

Zone [?](#) us-central1-c

Specify port name mapping (Optional)

Group type

Managed instance group
Managed instance group contains identical instances, created from an instance template, and supports autoscaling, autohealing, rolling updating, load balancing and more. VM instances are stateless and disks are deleted on VM deletion or recreation. [Learn more](#)

Unmanaged instance group
Unmanaged instance group is best for load balancing dissimilar instances, which you can add and remove arbitrarily. Autoscaling, autohealing, and rolling updating are not supported. [Learn more](#)

Instance template [?](#)





Setting Autoscaling settings

Autoscaling ?

On

Autoscale based on ?

For best results read [Configuring autoscaling instance groups](#)

CPU usage

Target CPU usage ?

Scaling dynamically creates or deletes VMs to meet the group target. [Learn more](#)

60

Cool-down period ?

60

seconds

Minimum number of instances ?

1

Maximum number of instances ?

10

Autohealing

VMs in the group are recreated as needed. You can use a health check to recreate a VM if the health check finds the VM unresponsive. If you do not select

[Create another health check](#)

No health check

elastic-cluster-health-check (TCP)

port: 9200, timeout: 5s, check interval: 5s, unhealthy threshold: 2 attempts

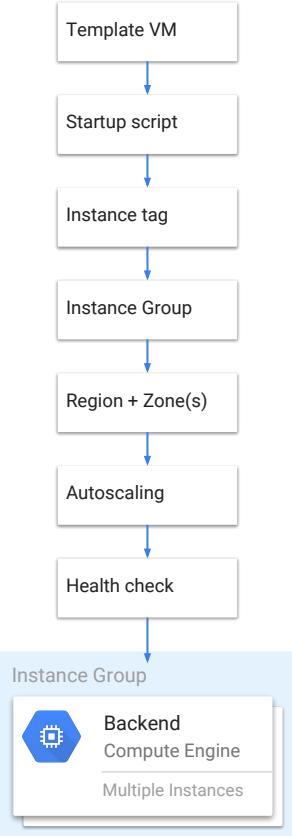
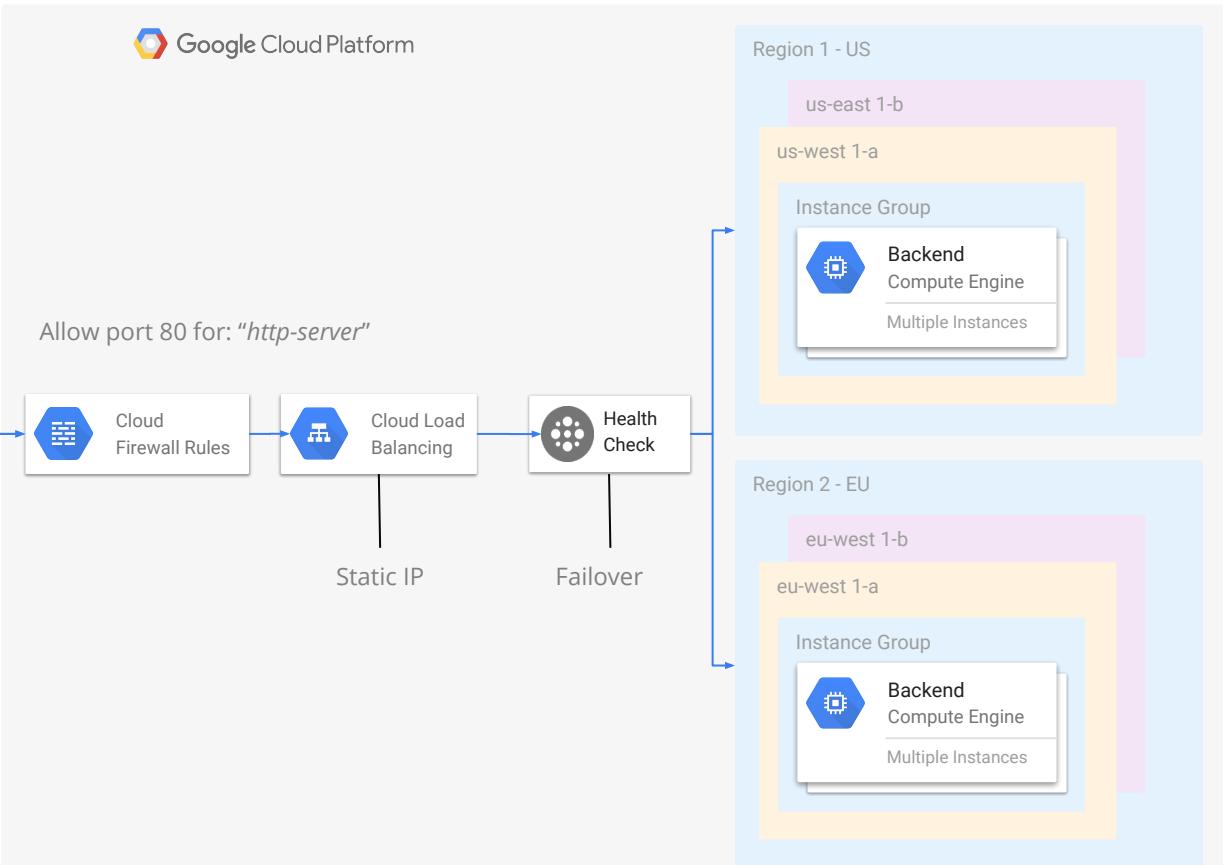
[Advanced creation options](#)

Create

Cancel



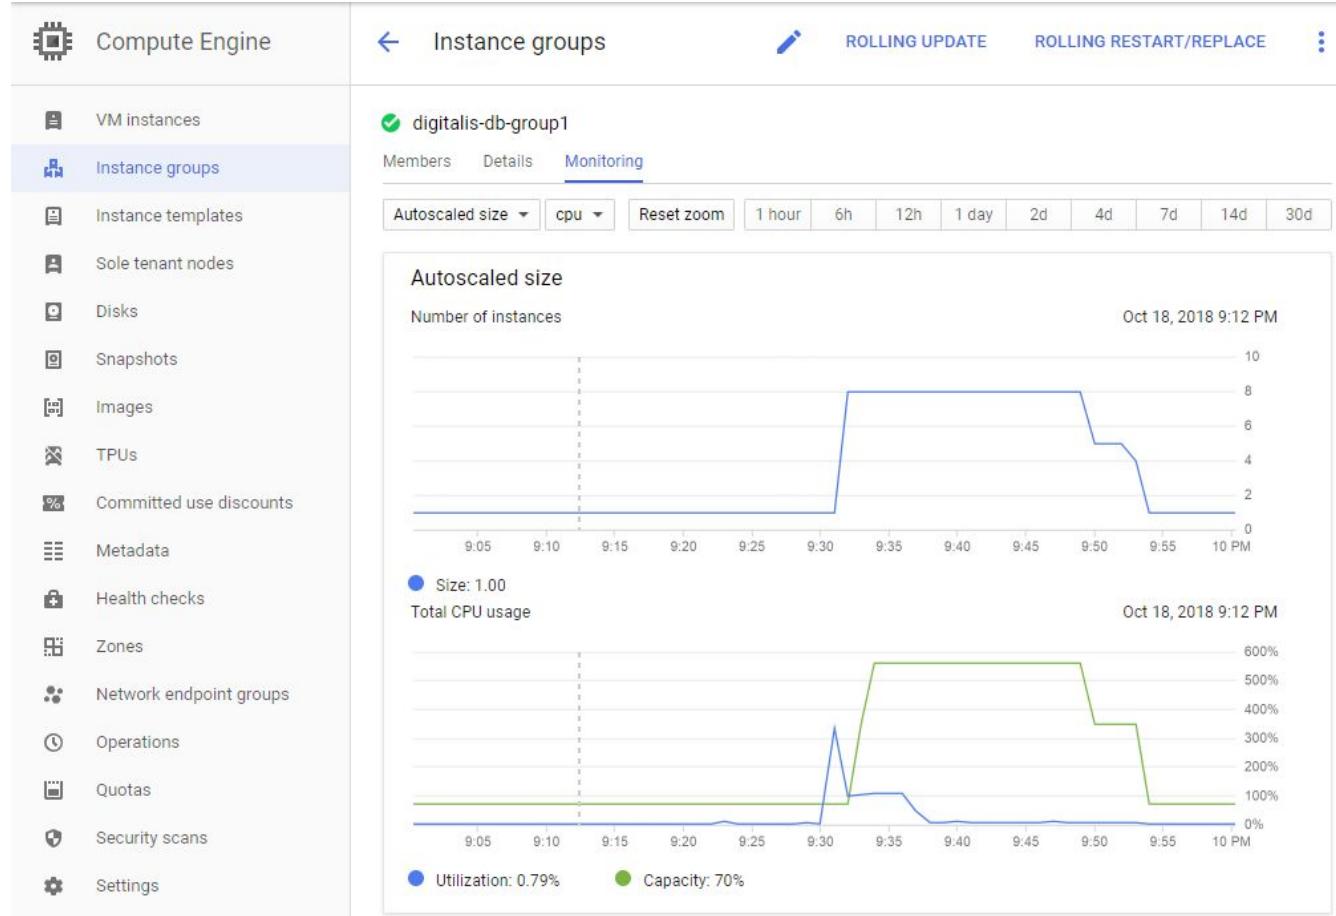
Instance Group - Autoscaling - Health Check



@martonkodok



Downscale on no load after 10 minute

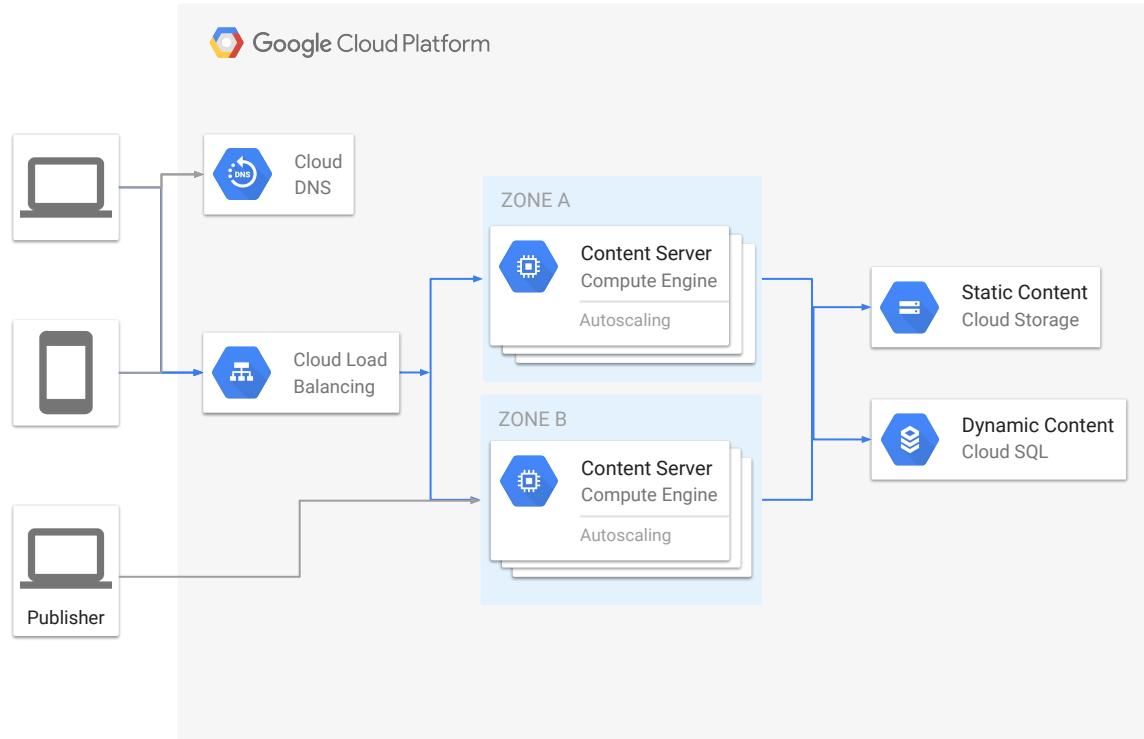




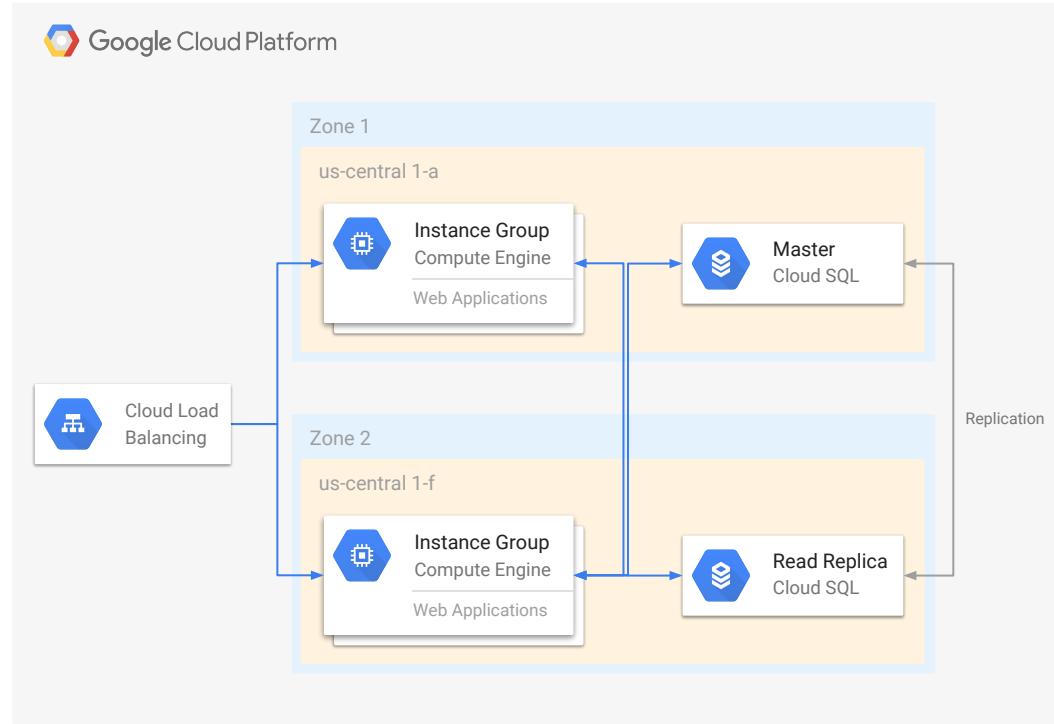
Example Architecture for *Autoscaling*



CMS with Autoscaling



DB multiple zones and Master/Slave topology

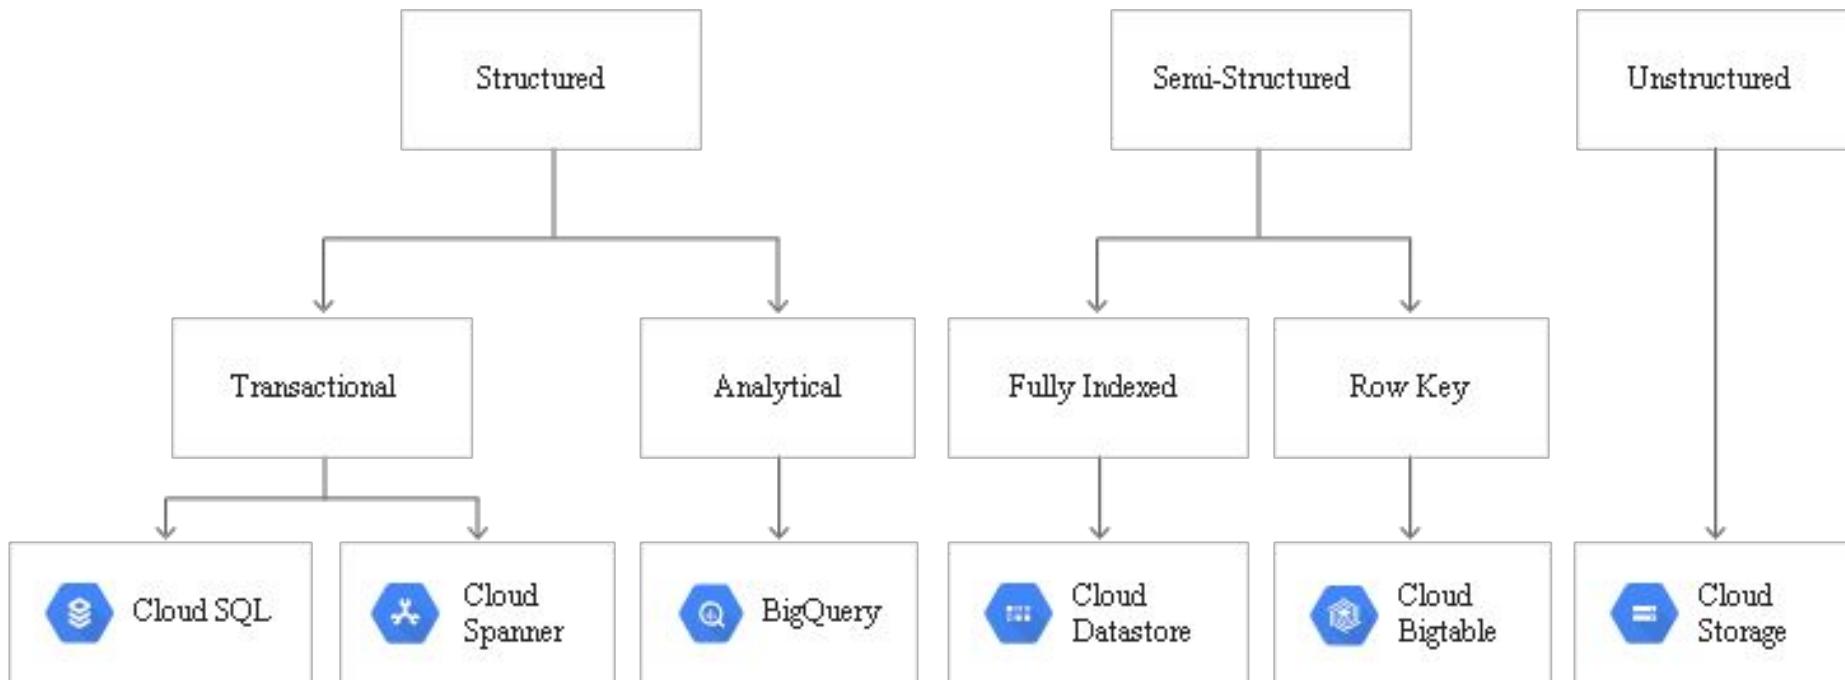




Data layer design



Database Mapping





Serverless



Meet Serverless



Automatic scaling



No server
management



Only pay for what
you use



Event-driven

Meet Serverless



Automatic scaling



No server
management



Only pay for what
you use



Event-driven



serverless data center depicted

Serverless is about maximizing **elasticity, cost savings, and agility** of cloud computing.



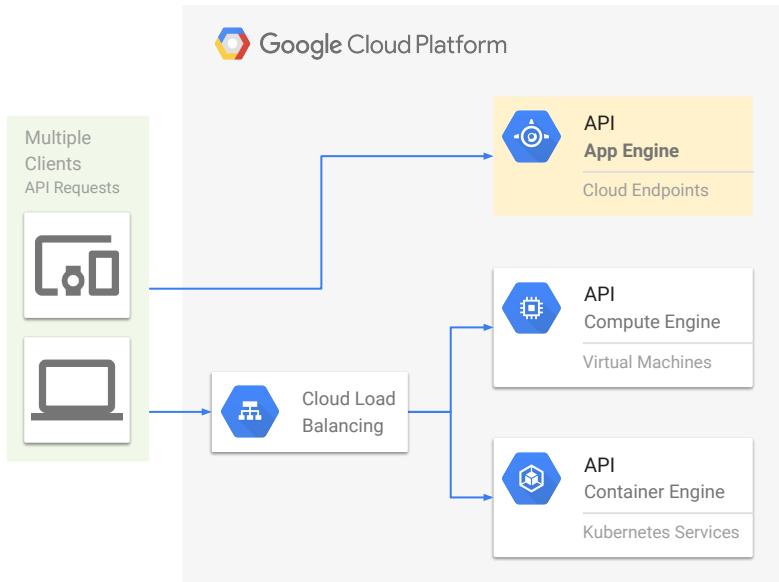
Serverless types

Platforms

Triggered Code



App Engine - application platform



App Engine

- Fully managed serverless application platform
- Scale your applications seamlessly from zero to planet scale
- Automatically scales depending on your application traffic
- Traffic Splitting (app versions, A/B tests, incremental rollouts)
- Standard vs Flexible environments

Python, Java, Node.js, Go, Ruby, PHP, .NET



App Engine

	Flexible	Standard
Instance startup time	Minutes	Seconds
Maximum request timeout	60 minutes	60 seconds
Background threads	Yes	Yes, with restrictions
SSH debugging	Yes	No
Writing to local disk	Yes, ephemeral (disk initialized on VM startup)	No
Modifying the runtime	Yes (through Dockerfile)	No
Supports installing third-party binaries	Yes	No
Pricing	Based on usage of vCPU, memory, and persistent disks	Based on instance hours

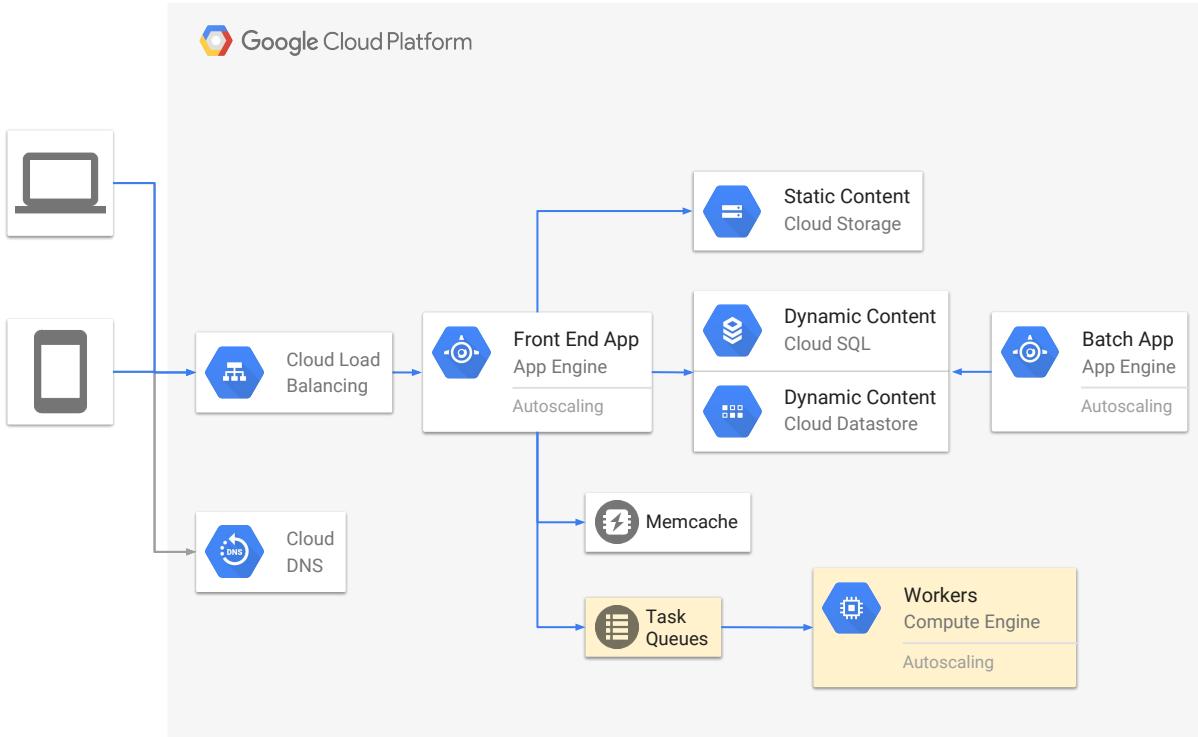




Message Queue systems: Task Queue Cloud Pub/Sub



Task Queue - Message Queue systems



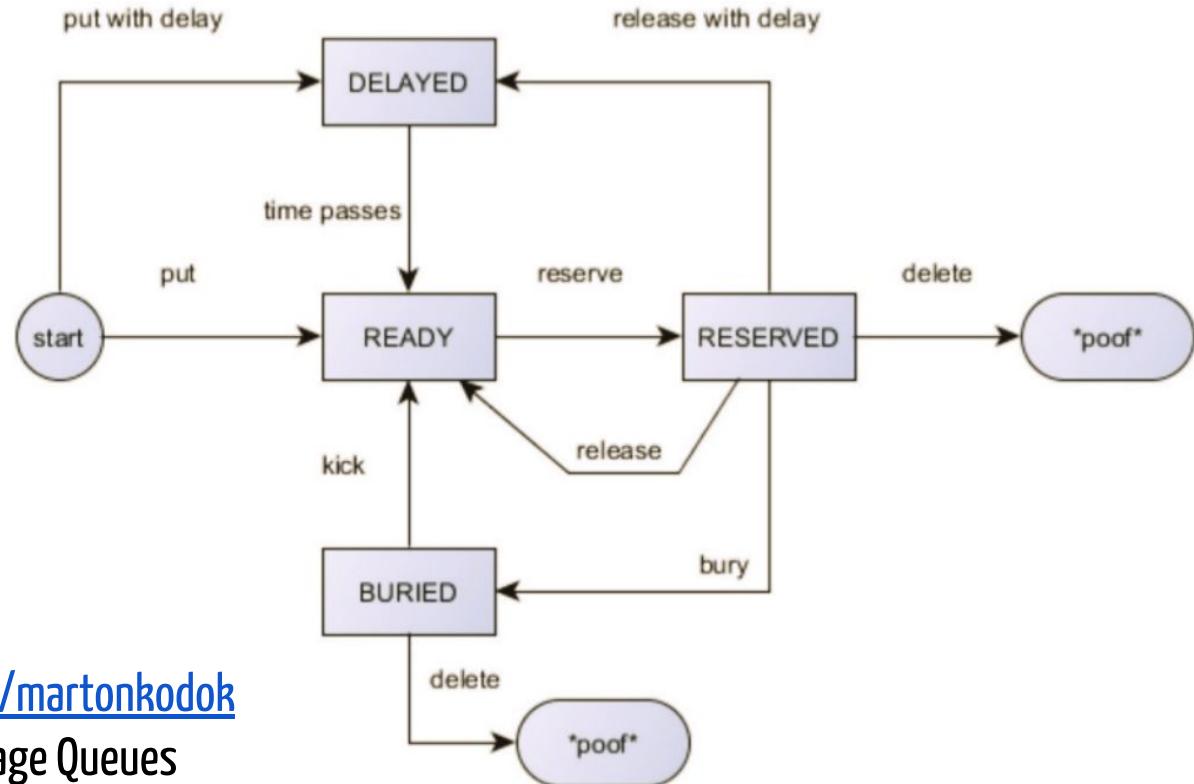
Task Queues

Push queues

- Put with delays

Pull queues

- Ability to “tag”
- Lease multiple eg: gameboard updated, game id as tag.



Slides: <https://www.slideshare.net/martondok>

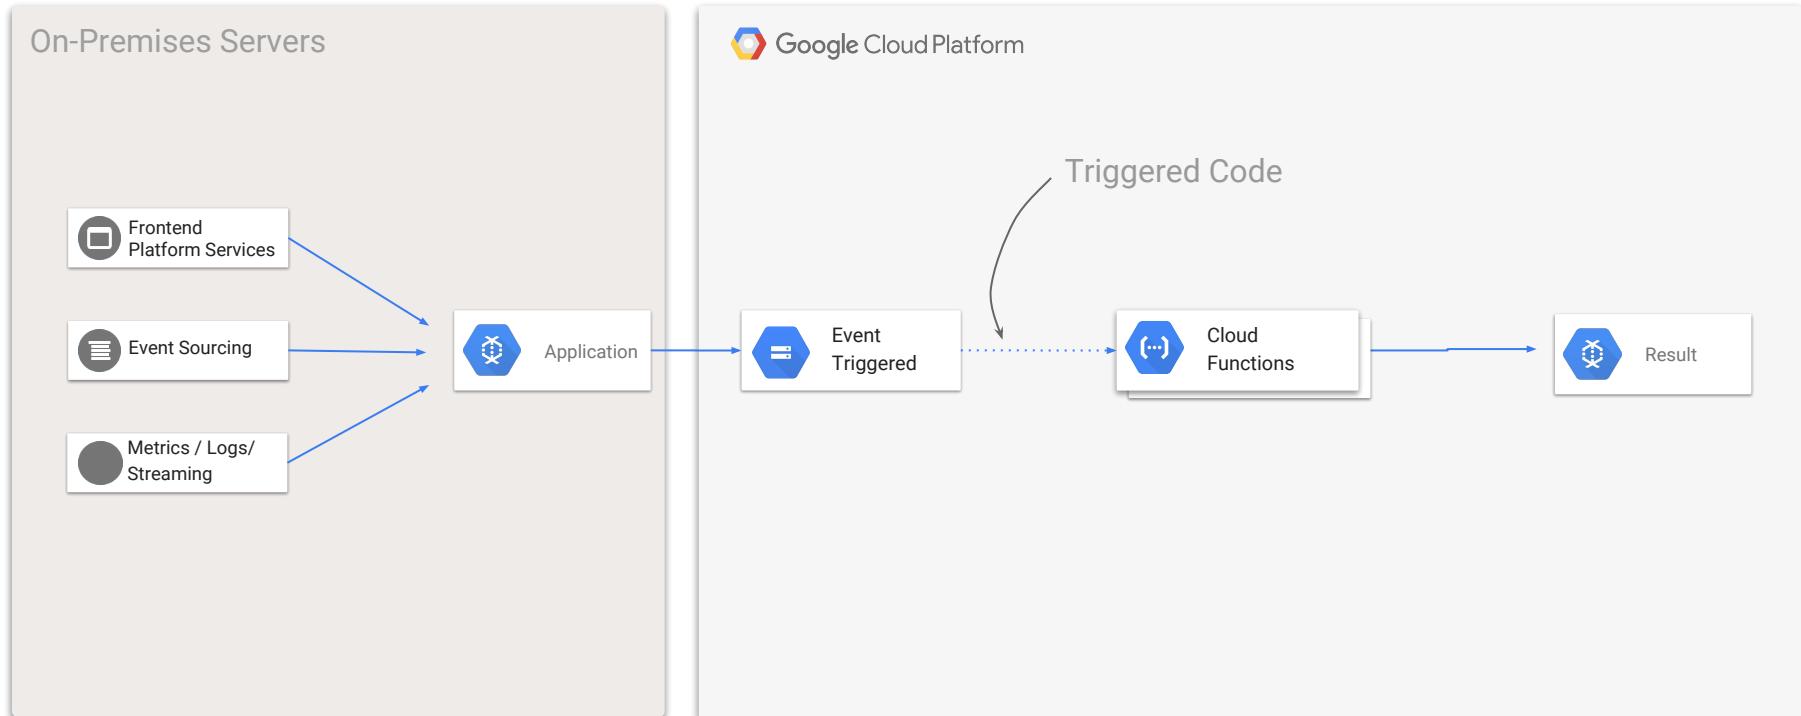
Title: Architectural Patterns - Message Queues



Cloud Functions



Cloud Functions



Cloud Functions - Event-driven - FaaS

- Simplest way to run your code in the cloud - just deploy
- Automatically scales, highly available and fault tolerant
- No servers to provision, manage, patch or update
- Pay only while your code runs
- Connects and extends cloud services

(In alpha: Go and Java)



- Node 8.11
- Support for async/await

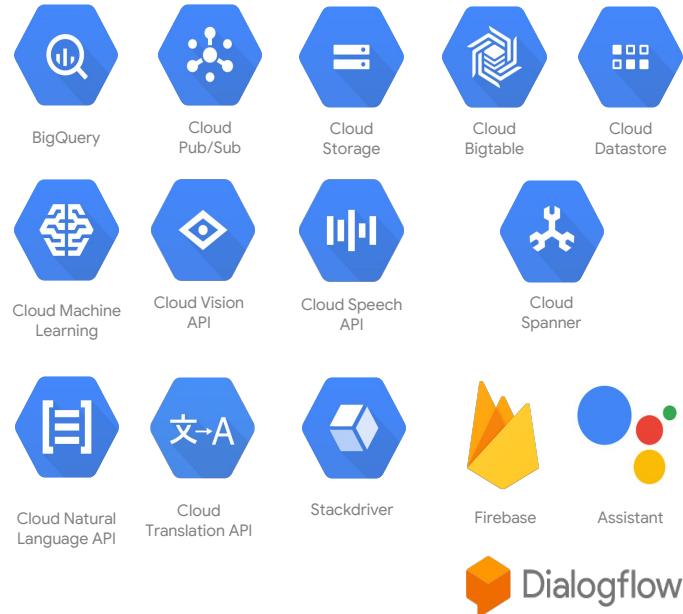
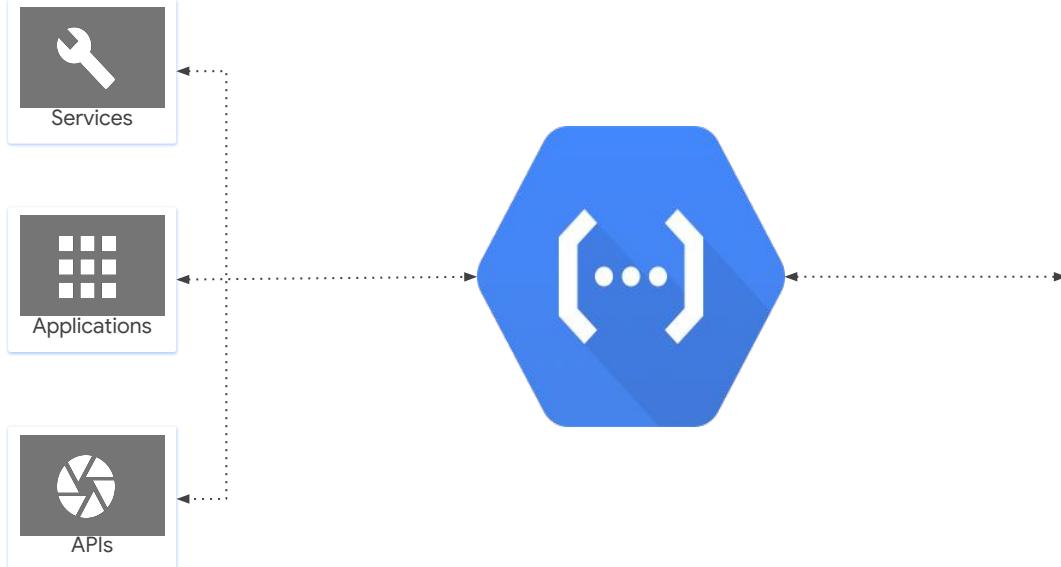


- Python 3.7.1
- Flask microframework



Functions: your gateway to GCP Services

Access 20+ Google services from GCF



Why is adopting serverless hard today?

1

Dependencies

Constrained runtimes,
frameworks and packages

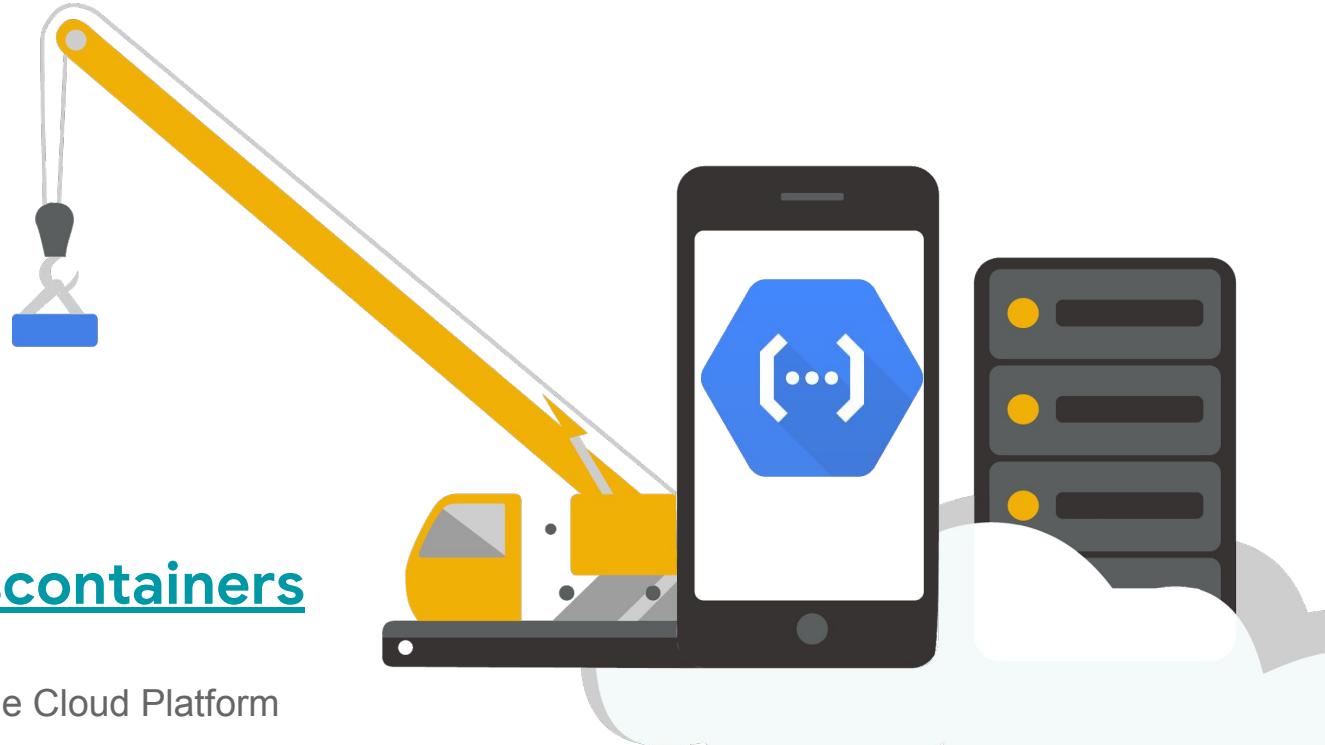
2

Multi-Cloud

Unable to run your workloads
on-prem, in the cloud or on a
third party service provider

New Serverless containers on GCF (EAP)

- Serverless containers
- Fully managed
- BYO workloads
- Pay for use
- Alpha invitations coming later this year



Sign up:
g.co/serverlesscontainers

Serverless means

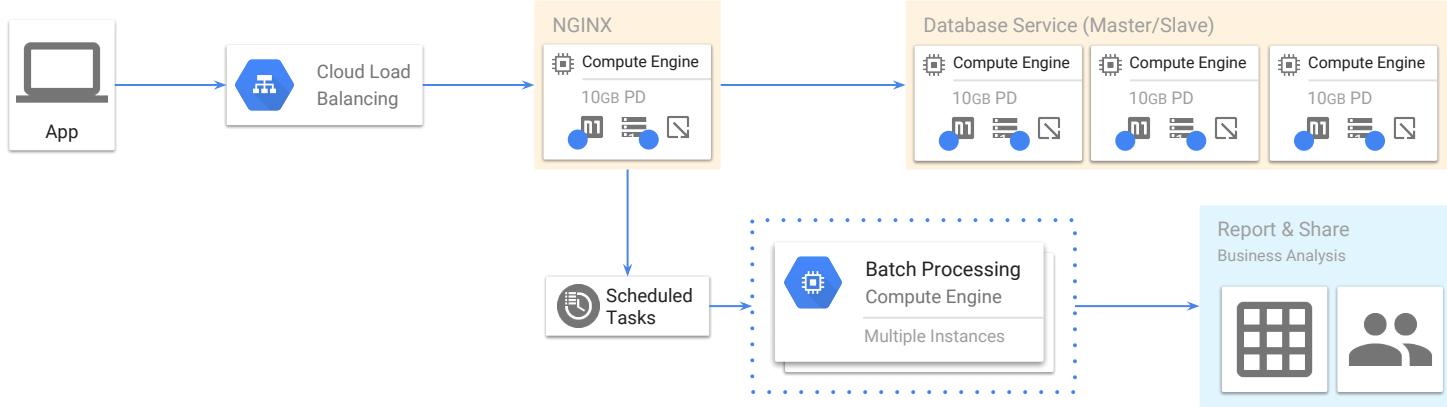
- No servers to provision or manage
- Scales with usage (ready every time for viral spikes or #BlackFriday)
- Availability and fault tolerance built in
- No orchestration in code
- Never pay for idle
- Decoupled: APIs as contracts (Abstract away the complexity)
- Monitored: Metrics and logging are a universal right
- Think concurrent, stateless, queue, stream based.



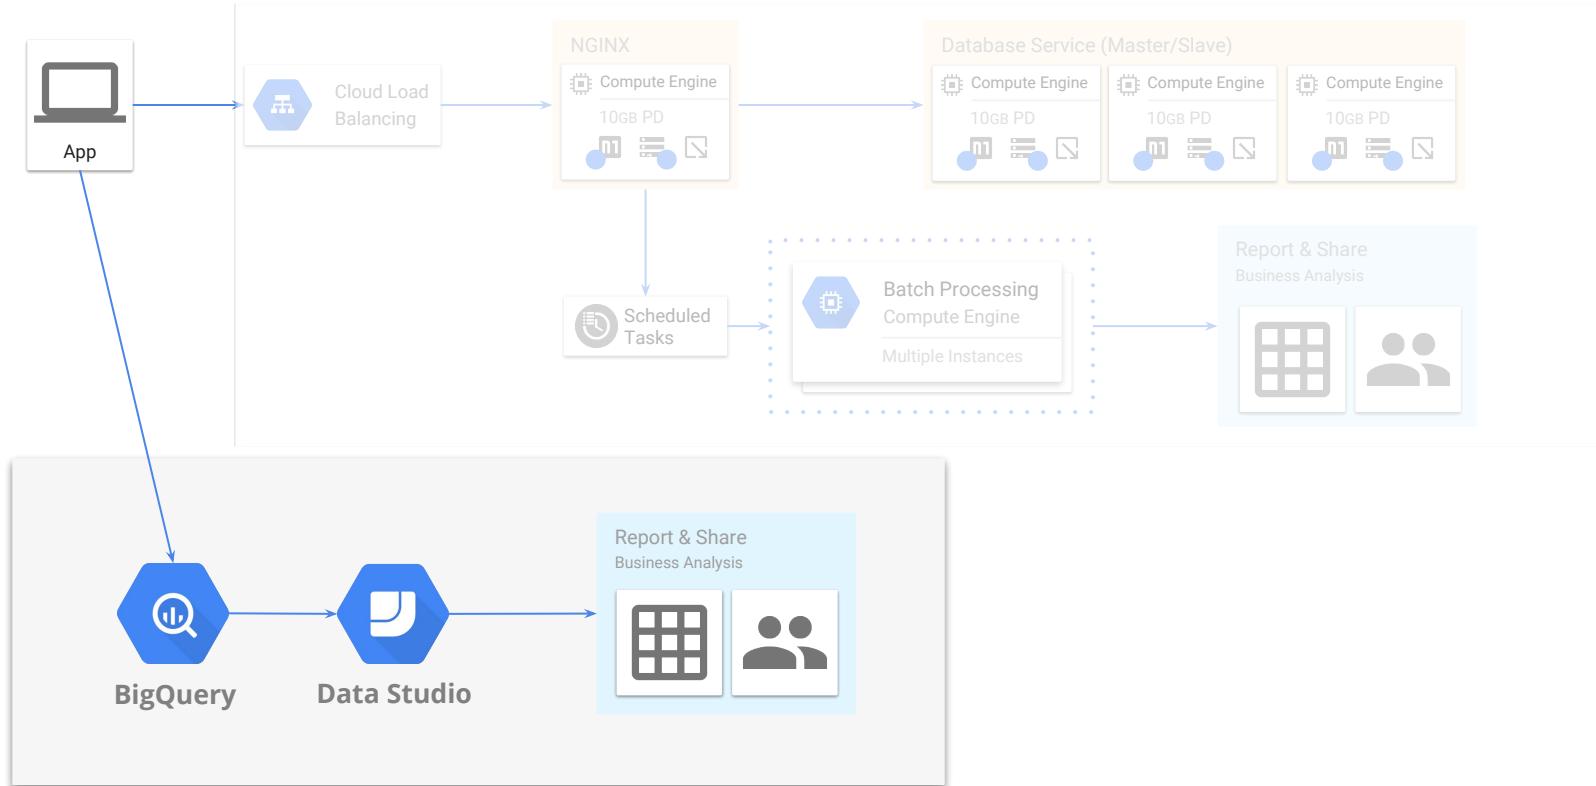
Reporting & Analytics moved to the Cloud



Legacy Reporting System



Serverless Reporting System



The Google BigQuery logo icon is a blue hexagon containing a white magnifying glass symbol with a bar chart inside it.

Google BigQuery



What is BigQuery?

Analytics-as-a-Service - Data Warehouse in the Cloud

Scales into Petabytes on Managed Google Infrastructure (US or EU zone)

SQL 2011 + Javascript UDF (User Defined Functions)

Familiar DB Structure (table, views, struct, nested, JSON)

Integrates with Google Sheets + Cloud Storage + Pub/Sub connectors

Decent pricing (queries \$5/TB, storage: \$20/TB cold: \$10/TB) *June 2018

Open Interfaces (Web UI, BQ command line tool, REST, ODBC)





BigQuery: Convenience of SQL

Columnar storage (max 10 000 columns in table)

Large files for loading: 5TB (CSV or JSON)

UDF in Javascript or SQL

Append-only tables preferred (DML syntax available)

Day column partitioned tables (select * from t where day='2018-01-01')

Rich SQL 2011: JSON, IP, Math, RegExp, Geocode, Window functions

Modern data types: Record, Nested, Struct, Array



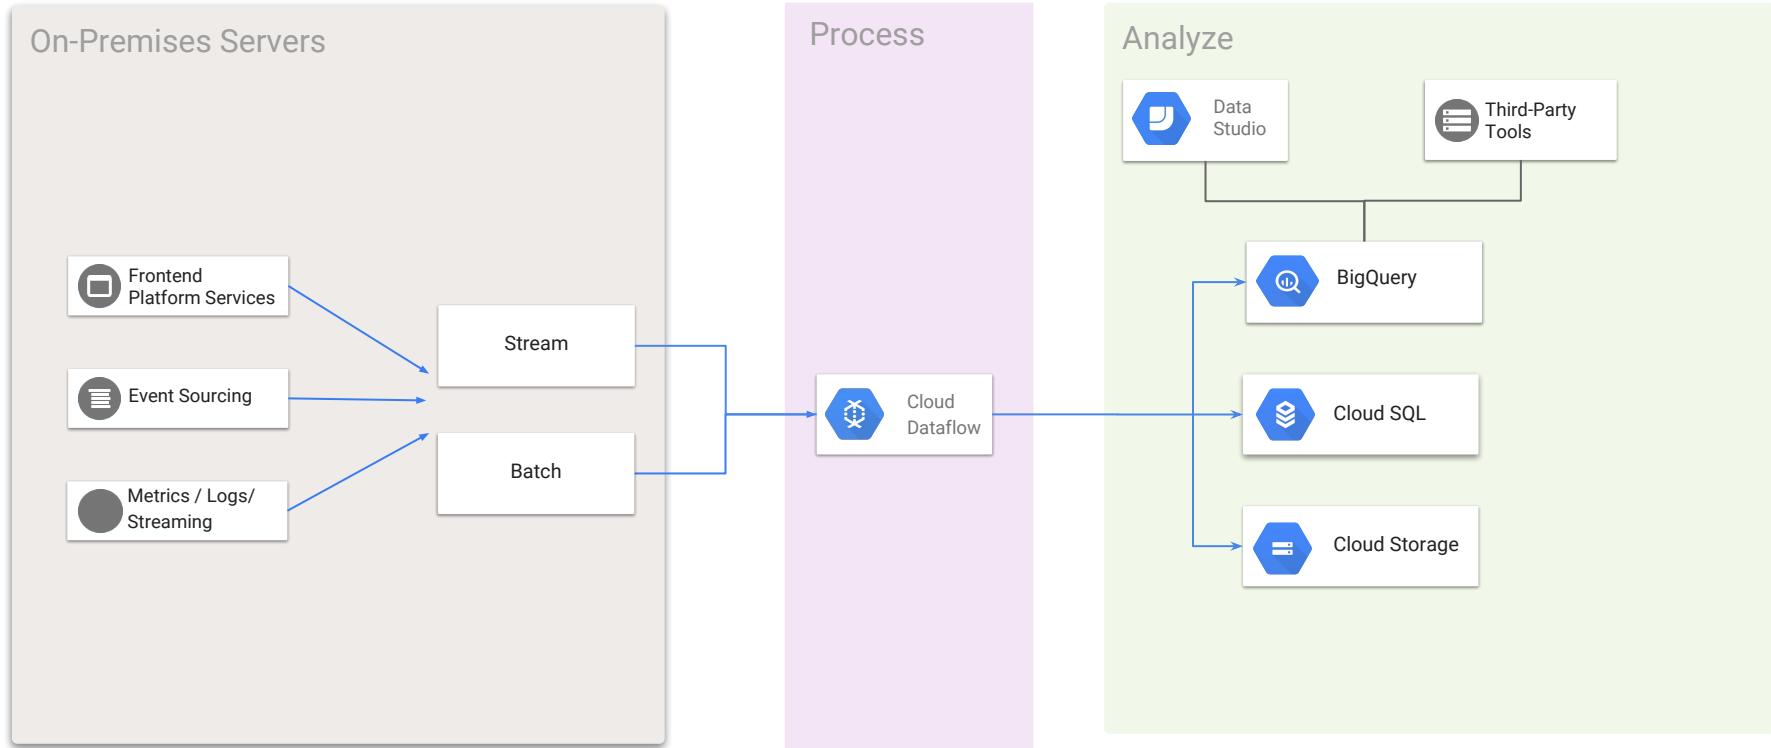
“



Data needs to be processed in
multiple services.
How can we pipe to multiple places?



Architecting for The Cloud



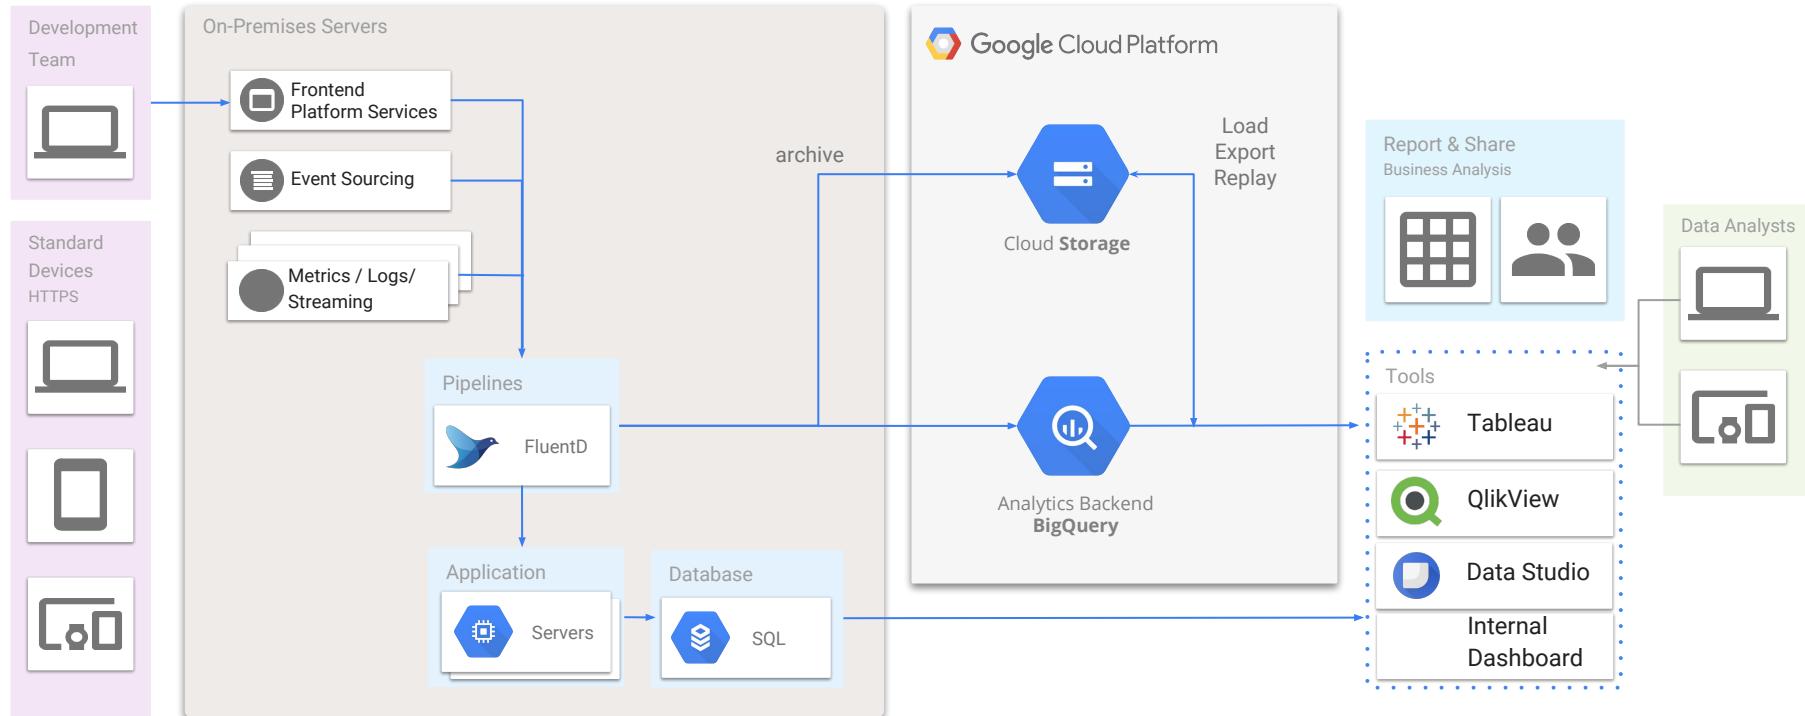
“



We have our app outside of GCP.
How can we use the benefits of BigQuery?



Data Pipeline Integration at REEA.net





Our benefits using BigQuery

- SQL language to run BigData queries
- run raw ad-hoc queries (either by analysts/sales or Devs)
- no more throwing away-, expiring-, aggregating old data
- no provisioning/deploy
- no running out of resources
- no more focus on large scale execution plan



Summary: When to use what



Cloud Functions

Smallest unit of computing

Event driven architecture

Connect & extend services



App Engine

Unit of computing is apps

HTTP request/response

Large scalable backends



Serverless add-on

Run functions, apps & containers on GKE

Full portability of your artifacts

Run on your own cluster

Thank you.

Slides available on: [slideshare.net/martonkodok](https://www.slideshare.net/martonkodok)

Reea.net - Integrated web solutions driven by creativity to deliver projects.

